

OmegaGo



A neural network trained 9x9 Go Playing Robot

Why brute force doesn't work



Possible positions:

9x9 board: 3^{81} (4.4×10^{38})

19x19 board: 3^{361} (1.7×10^{172})

Legal positions calculated by John Tromp

9x9 board: $\sim 1.039 \times 10^{38}$

19x19 board: $\sim 2.082 \times 10^{170}$

Atoms in the universe: 10^{80}


AlphaGo Techniques

- Neural net to predict top few moves (policy network)
- Does a Monte Carlo tree search on these top moves to determine the best one (with assistance of value network to prune the tree)


OmegaGo

- Strictly a neural network – no Monte Carlo tree search.
- OmegaGo has no knowledge of Go rules. If its first move choice is illegal it moves on to its 2nd choice.
- Trained on tens of thousands of 9x9 game records downloaded from the internet broken up into over 10 million moves.
- Games were played by players rated 5 kyu to pro

What is a neural network?



Zoomed in on one neuron



$$O = R\left(\sum_{j=1}^n w_j I_j + b\right)$$

TensorFlow



Feature

(Board position)

(Board position)

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & -1 & 1 & 1 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & -1 & 0 & 0 & 0 \\ 0 & 0 & -1 & 0 & 1 & -1 & 0 & -1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & -1 \\ 0 & 0 & 1 & 1 & 1 & 1 & 1 & -1 & 0 \\ 0 & 0 & 1 & -1 & 0 & -1 & -1 & 0 & -1 \\ 0 & 0 & 1 & -1 & 0 & 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Label
(Next Move)

(Next Move)

[illegible]

Feature

(Board position)

```
[ [ 0  0  0  0  0  0  0  0  0 ]
  [ 0 -1  1  1 -1  0  0  0  0 ]
  [ 0  0  0  1  0 -1  0  0  0 ]
  [ 0  0 -1  0  1 -1  0 -1  0 ]
  [ 0  0  0  0  0  0  0  1 -1 ]
  [ 0  0  1  1  1  1  1 -1  0 ]
  [ 0  0  1 -1  0 -1 -1  0 -1 ]
  [ 0  0  1 -1  0  0  0 -1  0 ]
  [ 0  0  0  0  0  0  0  0  0 ] ]
```

Label

(Next Move)

```
[ [ 0  1  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0 ] ]
```

Prediction

```
[ [ 0.    0.25  0.    0.    0.    0.    0.01  0.    0.    ]
  [ 0.    0.01  0.05  0.    0.    0.    0.01  0.    0.    ]
  [ 0.02  0.01  0.    0.02  0.    0.    0.01  0.    0.    ]
  [ 0.    0.    0.    0.01  0.    0.    0.    0.    0.    ]
  [ 0.    0.    0.01  0.02  0.01  0.01  0.2   0.01  0.    ]
  [ 0.    0.    0.    0.04  0.    0.01  0.    0.01  0.    ]
  [ 0.    0.    0.02  0.04  0.    0.    0.02  0.    0.    ]
  [ 0.01  0.03  0.    0.03  0.01  0.    0.    0.    0.07 ]
  [ 0.    0.    0.    0.    0.    0.    0.01  0.    0.    ] ]
```

Data Sources

Data Source	Number of Games	Ratings of Players
Top 50 players from app Go Quest	6 735	ELO rating above 2000 (equals about 1 Kyu)
Amateur games from app Go Quest	16 764	ELO ratings above 2000
Pro games from Mini-Go broadcasts by Yomiuri TV	413	Pro players
No Name Go Server (NNGS nngs.cosmic.org)	1 705	5 Kyu or better
Online-go.com	2 672	5 Kyu or better
TOTAL	28 289	
TOTAL # of moves	1 398 844	
TOTAL dataset after 8 transformations (flip, rotate board)	11 190 752	

SGF files: standard go format

```
(;  
GM[1]US[Brought to you by No Name Go Server]  
CoPyright[  
    This game was played on the No Name Go Server  
    Permission to reproduce this game is given.]  
GN[AdamSmith-vida(B) NNGS]  
EV[None]  
RE[W+2.5]  
PW[AdamSmith]WR[3k*]  
PB[vida]BR[5k*]  
PC[NNGS New York USA: ra.york.cuny.edu 9696]  
DT[1996-06-30]  
SZ[9]TM[300]KM[0.5]  
  
;B[fc];W[cg];B[cd];W[ff];B[gg];W[fg];B[gf];W[ge];B[he];  
W[gd];B[hd];W[gc];B[fd];W[gh];B[hh];W[fe];B[fh];W[eh];  
B[gi];W[dc];B[dd];W[ec];B[ed];W[fb];B[cc];W[cb];B[db];  
W[eb];B[bb];W[da];B[bf];W[ba];B[ac];W[bg];B[cf];W[af];  
B[ae];W[ag];B[bd];W[hc];B[dh];W[dg];B[ei];W[eg];B[di];  
W[df];B[de];W[id];B[if];W[ch];B[ci];W[bi];B[fi];W[ee];  
B[ie];W[ic];B[ab];W[ca];B[tt];W[ig];B[hf];W[tt];B[tt]  
;)
```

Training Results

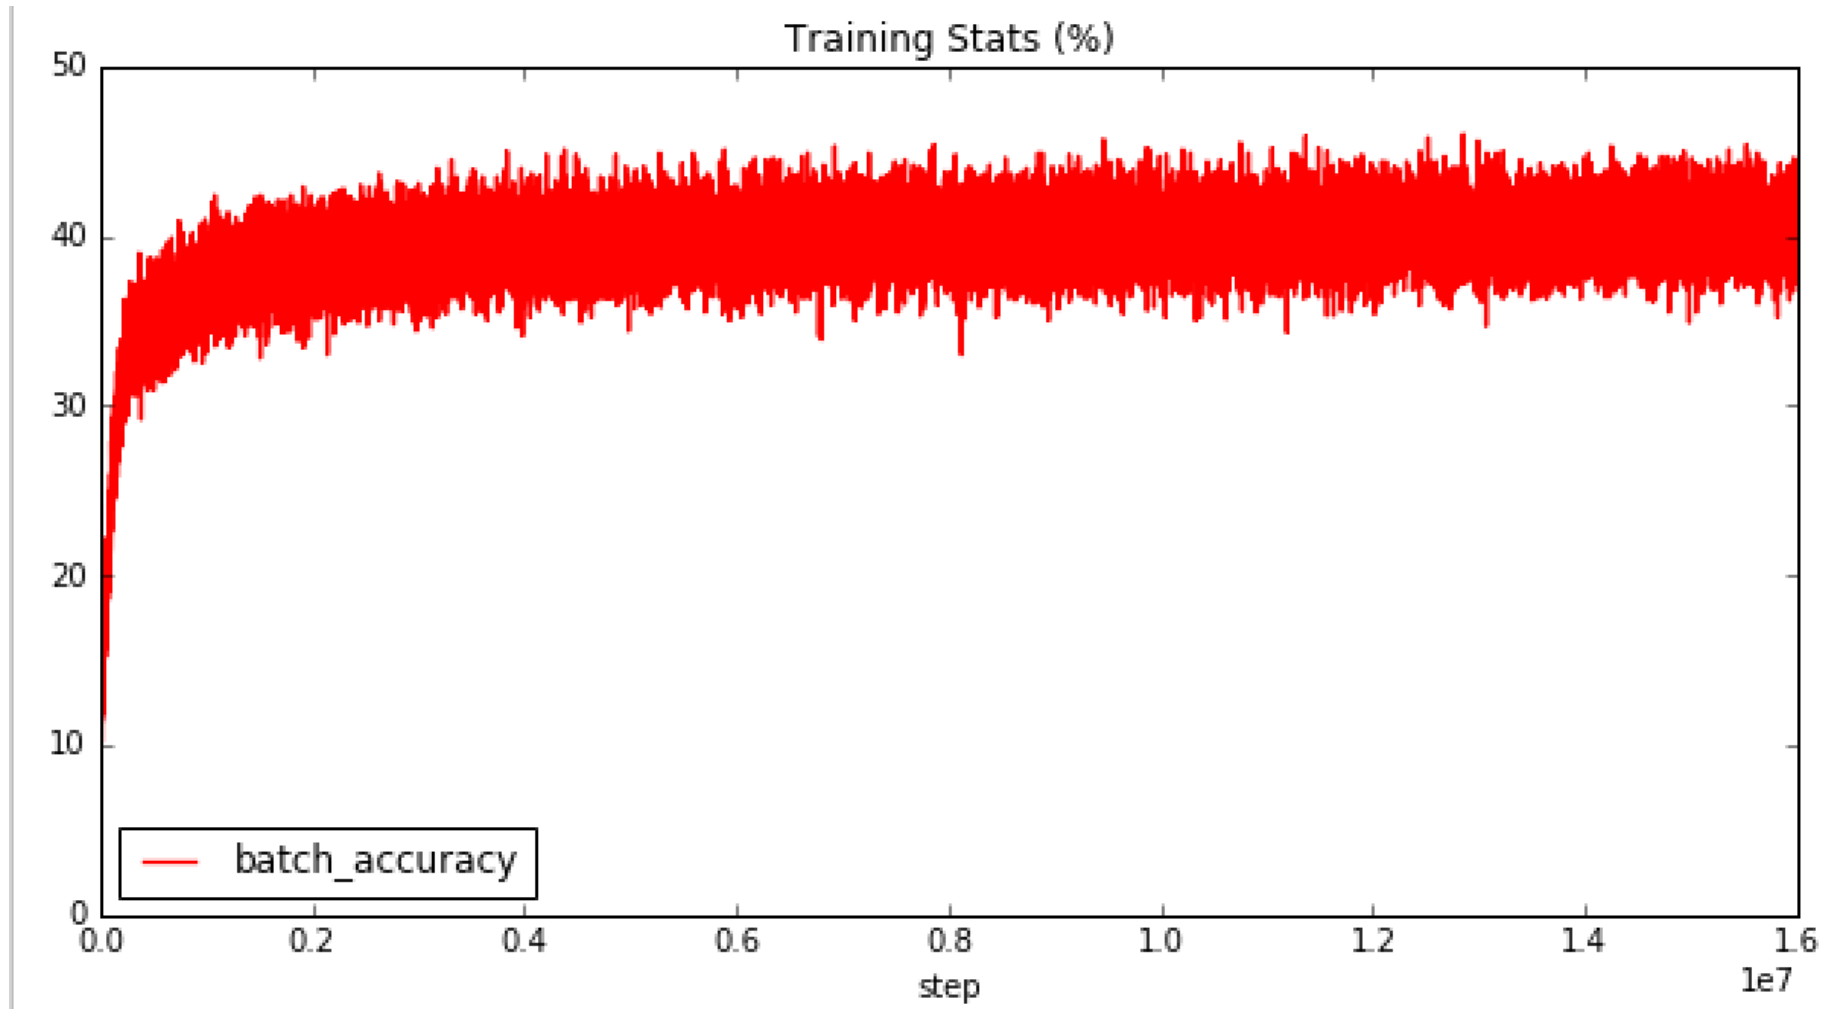
Test accuracy: 37%

AlphaGo's accuracy:
In the 50-60% range



Training Results 500,000 runs



Training Results 16 million runs



Prediction Array = Heat Map



```
for index in range(len(predictionArray)):
```

```
    x = index % 9
```

```
    y = index / 9
```

```
    heatEdgeColor = (.2,.2, .2)
```


```
    heatFaceColor = 'b'
```

```
    prediction = 200 * predictionArray[index]
```

```
# draw heat map
```

```
s2, = ax.plot(x,y,'s',markersize=prediction,  
              markeredgecolor=heatEdgeColor,  
              markerfacecolor=heatFaceColor,  
              markeredgewidth=2, alpha=0.3)
```

Handling Pass and Resign



```
if c.atEnd:
    # Add in the last board position
    # - the final pass or resign move.
    # The robot needs to train
    #when to pass and when to resign

    if (gameEndedByTime == True):
        break
        #No passing move needed

    label = np.zeros(83)

    if (gameEndedInResignation == True):
        # Board has just been flipped
        label[82]=1
    else:
        label[81]=1
```

Playing against GnuGo

A free software program by the Free Software Foundation that plays Go using the Go Text Protocol (GTP)



```
1 boardsize 9
=1
2 play black D5
=2
3 genmove white
=3 C3
4 play black C3
?4 illegal move
```

Demo of OmegaGo Versus GnuGo

GnuGo versus OmegaGo Stats

Even Game

(OmegaGo plays black.)

Komi is 6.5)

48% win

One stone handicap

(OmegaGo plays black.)

Komi is 0.5)

61% win

Even Game

(OmegaGo plays white.)

Komi is 6.5)

53% win

Two stone handicap

(OmegaGo plays black.)

Komi is 0.5)

85% win

Playing on KGS (NubbyBot)

Use kgsGtp.jar to connect

KGS: User "NubbyBot"

User Data Games Tagged Games Rank

Name: NubbyBot Email: Private

Real Name: Alexandra Patz


Rank: - Client version: English (United States)

Registered User Last On: Now

Registered On: Dec 22, 2016 ☒ Buddy? ☐ Censored? ☐ Fan?

Notes:

This is a neural net trained bot that plays at around the 20kyu level on 9x9 boards only.



OK

KGS Results

KGS: User "N"


User DataGamesTagged GamesRank

Games:

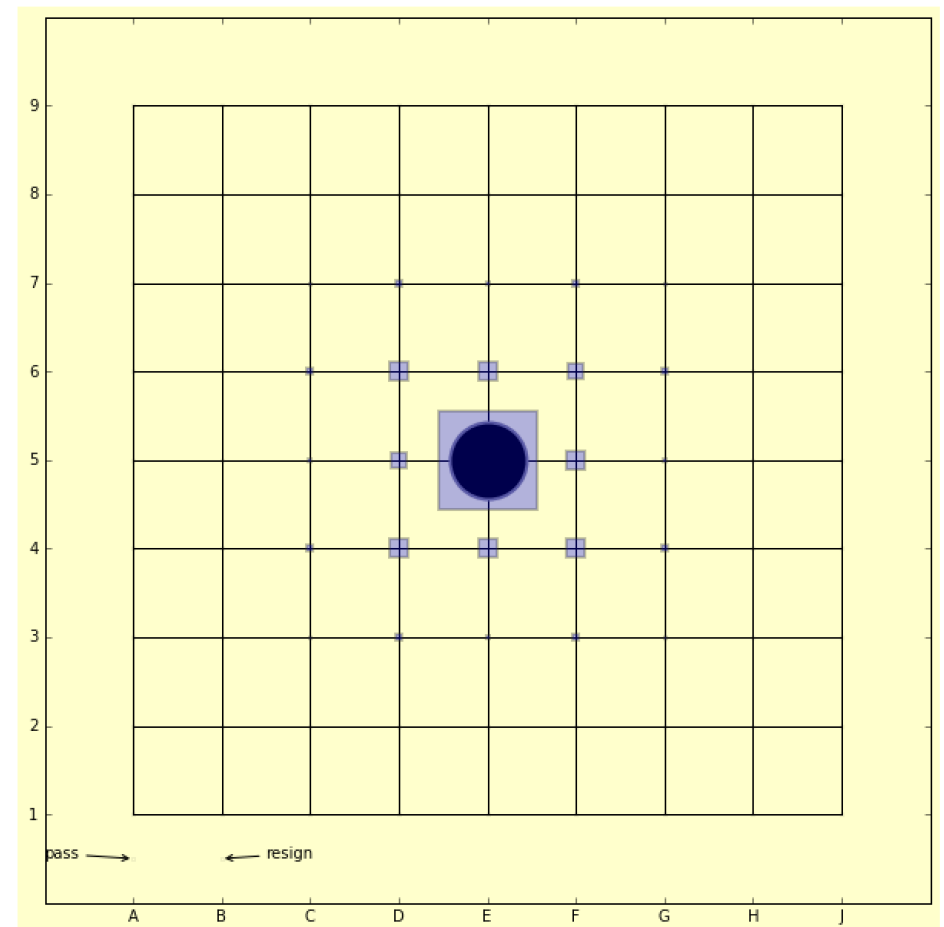
F	targogo [22k]	●NubbyBot [-]	9×9	Km 25.5	B+4.5	1/20/17 10:15:43 AM
F	grrrr [15k?]	●NubbyBot [-]	9×9	Km 6.5	B+7.5	1/20/17 10:04:24 AM
F	●NubbyBot [-]	grrrr [15k?]	9×9	Km 6.5	W+15.5	1/20/17 10:03:02 AM
F	●NubbyBot [-]	grrrr [15k?]	9×9	Km 6.5	W+12.5	1/20/17 10:01:22 AM
F	●NubbyBot [-]	grrrr [15k?]	9×9	Km 6.5	W+6.5	1/20/17 9:59:45 AM
F	●Nogosai [22k]	NubbyBot [-]	9×9	Km 6.5	W+37.5	1/20/17 9:56:55 AM
F	grrrr [15k?]	●NubbyBot [-]	9×9	Km 6.5	B+4.5	1/20/17 9:54:17 AM
F	●NubbyBot [-]	grrrr [15k?]	9×9	Km 6.5	W+2.5	1/20/17 9:52:23 AM
F	●NubbyBot [-]	grrrr [15k?]	9×9	Km 6.5	W+21.5	1/20/17 9:50:30 AM
F	NubbyBot [-]	●tengen86 [4k]	9×9	Km 6.5	B+24.5	1/20/17 9:48:52 AM
F	●NubbyBot [-]	dudi11	9×9	Km 6.5	W+5.5	1/20/17 9:44:44 AM
F	NubbyBot [-]	xcoi [?]	19×19	Km 6.5		1/20/17 9:20:45 AM
F	tuananh	●NubbyBot [-]	9×9	Km 6.5	B+94.5	1/20/17 9:13:15 AM
F	●NubbyBot [-]	upgrade2	9×9	Km 6.5	W+3.5	1/20/17 9:07:46 AM
F	NubbyBot [-]	●murx	9×9	Km 6.5	B+15.5	1/20/17 9:01:04 AM
F	●dabe005 [-]	NubbyBot [-]	9×9	Km 6.5	W+17.5	1/20/17 8:57:00 AM
F	●NubbyBot [-]	sanfran [-]	9×9	Km 6.5	W+Res.	1/20/17 8:54:20 AM
F	NubbyBot [-]	rmk [?]	9×9	Km 6.5		1/20/17 8:31:15 AM
F	rmk [?]	NubbyBot [-]	9×9	Km 6.5		1/20/17 8:11:15 AM
F	rmk [?]	NubbyBot [-]	9×9	Km 6.5		1/20/17 8:01:54 AM
F	NubbyBot [-]	●umayado [?]	9×9	Km 6.5	B+9.5	1/20/17 7:55:16 AM
F	●NubbyBot [-]	umayado [?]	9×9	Km 6.5	W+Res.	1/20/17 7:53:37 AM
F	umayado [?]	●NubbyBot [-]	9×9	Km 6.5	B+Res.	1/20/17 7:51:04 AM
F	NubbyBot [-]	●umayado [?]	9×9	Km 6.5	B+3.5	1/20/17 7:46:09 AM
F	●NubbyBot [-]	deorde [19k]	9×9	Km 6.5	W+22.5	1/20/17 7:42:41 AM
F	●NubbyBot [-]	deorde [19k]	9×9	Km 6.5	W+17.5	1/20/17 7:39:03 AM
F	NubbyBot [-]	●DeathM38 [?]	9×9	Km 6.5	B+74.5	1/20/17 7:30:05 AM

Human Bias









No data transformation: Note human bias



Eight transformations performed



Transforming data to increase data size and overcome human bias



Using `numpy.fliplr()`, `numpy.flipud()` and `numpy.rot90()`

Neural Network Design

Dataset: 11 million features from strong games

Optimizer: GradientDescentOptimizer

Learning rate: 0.05

Layers: 2 (one hidden layer)

Activation function: Relu

Hidden layer number of nodes: 1024

Batch size: 1000

Runs: 20 million

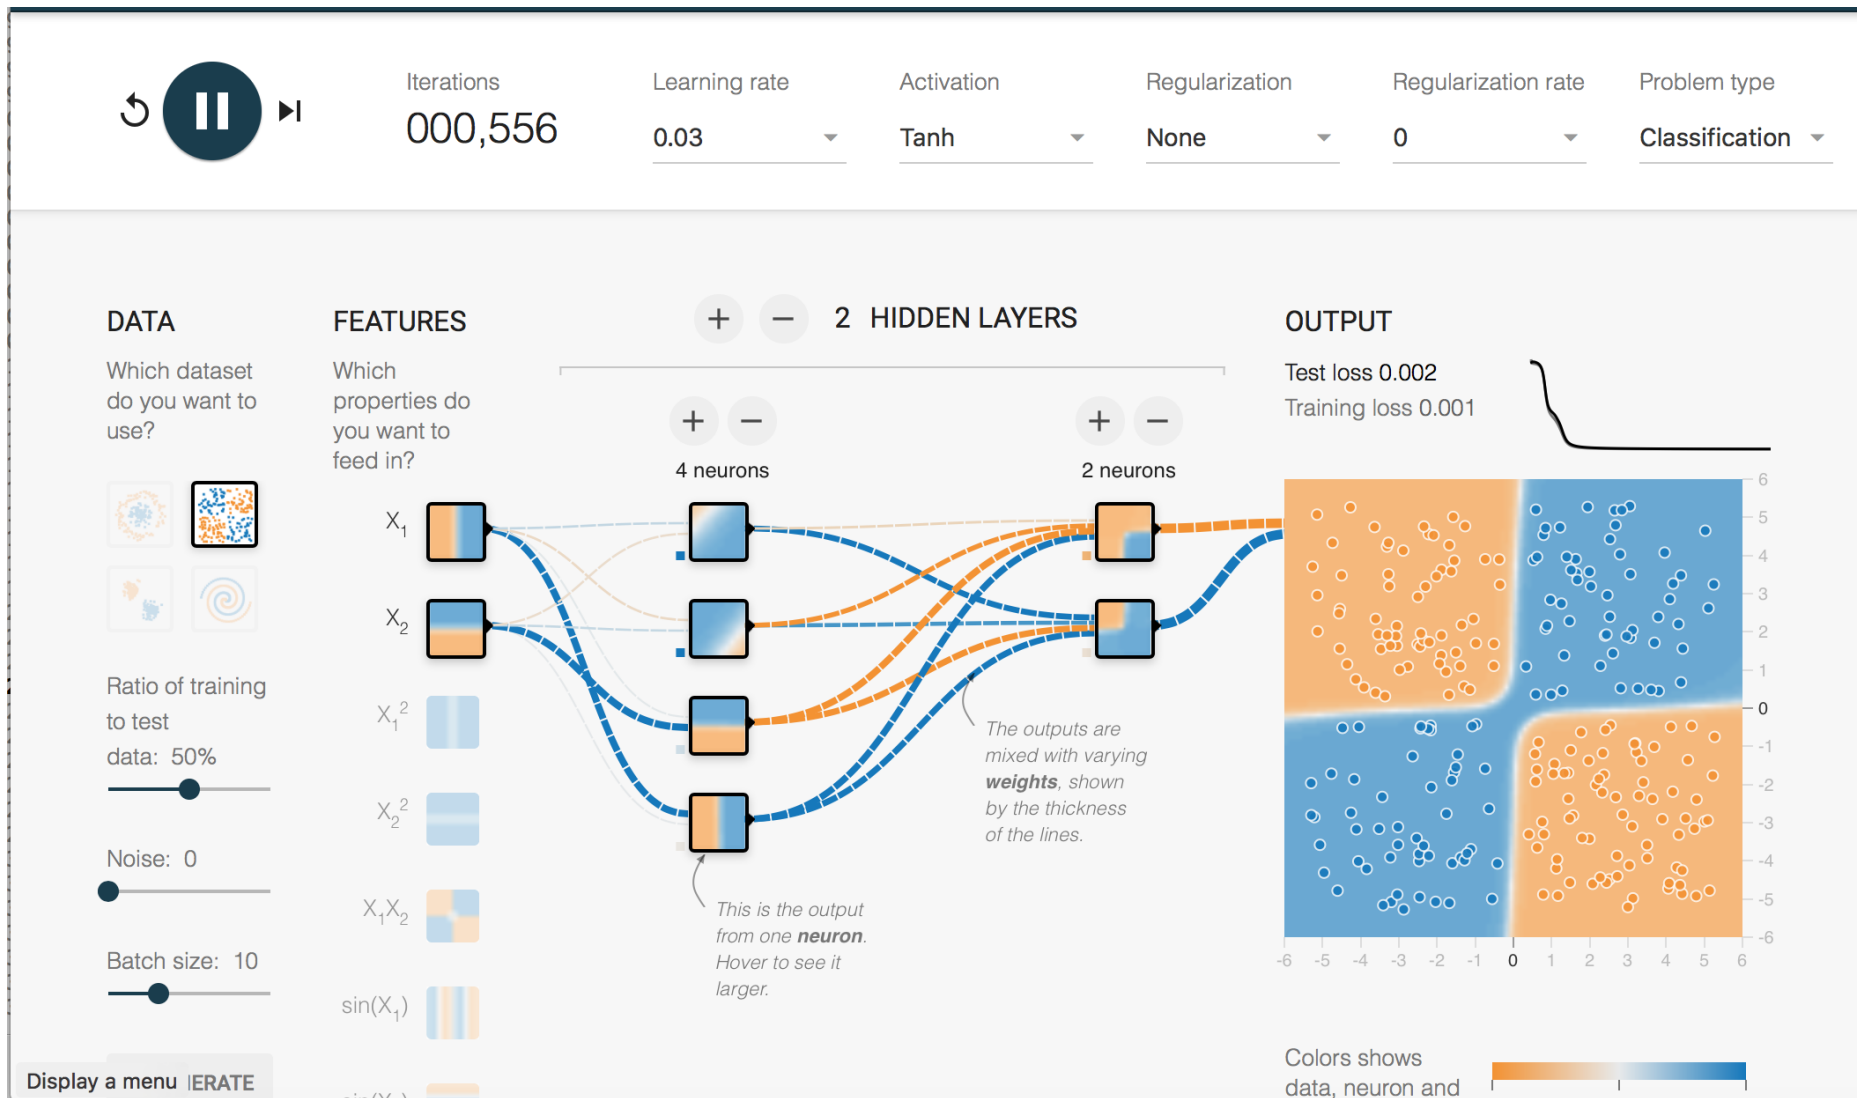
Shuffling after each full pass through the dataset

Hyper parameter variations tried:

- Three layer network (2 hidden layers)
- More hidden nodes / less hidden nodes
- Different optimizer algorithm (AdagradOptimizer)
- Different activation function (tanh instead of relu)
- Decaying learning rate

TensorFlow Playground

playground.tensorflow.org



Other variations tried

- Added mixed level games to the dataset
Black: >16 kyu
White: < 6 kyu
Improved win rates against gnuGo by 5%!
- Tried adding another feature plane

Adding another feature plane


RESEARCH ARTICLE

Extended Data Table 2 | Input features for neural networks

Feature	# of planes	Description
Stone colour	3	Player stone / opponent stone / empty
Ones	1	A constant plane filled with 1
Turns since	8	How many turns since a move was played
Liberties	8	Number of liberties (empty adjacent points)
Capture size	8	How many opponent stones would be captured
Self-atari size	8	How many of own stones would be captured
Liberties after move	8	Number of liberties after this move is played
Ladder capture	1	Whether a move at this point is a successful ladder capture
Ladder escape	1	Whether a move at this point is a successful ladder escape
Sensibleness	1	Whether a move is legal and does not fill its own eyes
Zeros	1	A constant plane filled with 0
Player color	1	Whether current player is black

Feature planes used by the policy network (all but last feature) and value network (all features).

More Feature Planes: Legal and Illegal Moves



One-hot-encoded feature planes

[0. 0. 0. 0. 0. 0. 0. 0. 0.
 1. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 1. 0. 1. 0. 0. 0.
 0. 0. 0. 0. 1. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 1. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]

Black

0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 1. 0. 0. 0. 0.
 0. 0. 0. 1. 0. 1. 0. 0. 0.
 0. 0. 0. 0. 1. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 1.
 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]

White

1. 1. 1. 1. 1. 1. 1. 1. 1.
 0. 1. 1. 1. 0. 1. 1. 1. 1.
 1. 1. 1. 0. 0. 0. 1. 1. 1.
 1. 1. 1. 0. 0. 0. 1. 1. 1.
 1. 1. 1. 1. 0. 1. 1. 1. 1.
 1. 1. 1. 1. 1. 1. 1. 1. 1.
 1. 1. 1. 1. 1. 1. 1. 1. 1.
 0. 1. 1. 1. 1. 1. 1. 1. 0.
 1. 0. 1. 1. 1. 1. 1. 0. 0. 0. 0.]

**Empty
Legal**

0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 1. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]

**Empty
Illegal**

Code on GitHub:

https://github.com/nubbyp/omega_go